

UNIT I: Introduction to Deep Learning

Q1. Compare and contrast the working of the human brain with that of a neural network.

Aspect	Human Brain	Neural Network
Basic Unit	Neuron (biological)	Artificial Neuron
Structure	~86 billion neurons interconnected	Layers of artificial neurons
Signal Transmission	Uses electrochemical signals across synapses	Uses numerical values and weights across nodes
Learning Mechanism	Learns through synaptic plasticity and neurochemical changes	Learns through adjusting weights using algorithms like backpropagation
Processing	Highly parallel and distributed	Simulated parallelism using hardware (GPUs/TPUs)
Adaptability	Highly adaptable and generalizes well in unseen situations	Needs large data and training; generalization is limited by design
Energy Efficiency	Very efficient (~20W for the entire brain)	High power consumption, especially during training

Artificial Neural Networks (ANNs) are inspired by the biological brain but are simplified mathematical models. While ANNs mimic the learning process of the brain, they are limited in complexity and adaptability compared to the human brain.

Q2. Describe various types of data handled by deep learning models with examples.

Deep learning models are designed to handle a wide variety of data types. Each type requires different preprocessing and model architecture.

1. Time Series Data:

- **Definition:** Data collected sequentially over time.
- **Examples:** Stock prices, weather data, ECG signals, sensor readings.
- **Model Used:** Recurrent Neural Networks (RNNs), Long Short-Term Memory (LSTM), GRUs.
- **Use Case:** Forecasting future values (e.g., predicting next day's temperature or stock prices).

2. Image Data:

- **Definition:** Data in the form of pixel arrays, usually 2D for grayscale or 3D for RGB images.
- **Examples:** Photographs, X-rays, satellite imagery.
- **Model Used:** Convolutional Neural Networks (CNNs).
- **Use Case:** Object detection, facial recognition, medical imaging diagnostics.

3. Video Data:

- **Definition:** Sequence of image frames over time (i.e., a combination of image and time series data).
- **Examples:** Surveillance videos, YouTube clips, autonomous driving footage.
- **Model Used:** 3D CNNs, CNN + RNN hybrid models.
- **Use Case:** Activity recognition, video summarization, gesture detection.

Deep learning models can be tailored to handle structured and unstructured data types like time series, images, and videos by choosing the appropriate architecture and preprocessing techniques.

Q3. Explain the common architectural principles followed in designing deep learning models.

The architecture of deep learning models is designed using well-defined principles to ensure that they can effectively learn from complex and high-dimensional data. Below are the major architectural principles:

1. Layered Structure (Input–Hidden–Output):

- Deep learning models are composed of layers of interconnected neurons.
- The **input layer** receives raw data (e.g., image pixels or feature vectors).
- One or more **hidden layers** perform feature extraction and transformations.
- The **output layer** generates final predictions, such as a class label or numerical value.

2. Depth and Width of the Network:

- **Depth** refers to the number of hidden layers; deeper networks capture higher-level features.
- **Width** refers to the number of neurons in each layer.
- Example: In a CNN, early layers detect edges, and deeper layers detect shapes and objects.

3. Activation Functions:

- These introduce **non-linearity**, allowing the model to learn complex mappings.
- **ReLU (Rectified Linear Unit)**: Speeds up training and reduces vanishing gradient issues.
- **Sigmoid and Tanh**: Used in shallow networks or specific layers (e.g., output).
- Without activation functions, a neural network behaves like a linear model.

4. Weight Initialization:

- Poor initialization can cause slow convergence or gradient issues.
- **Xavier initialization** (for sigmoid/tanh) and **He initialization** (for ReLU) help maintain stable gradients throughout the network.

5. Regularization Techniques:

- Prevents overfitting and improves generalization to unseen data.
- **Dropout**: Randomly deactivates neurons during training.
- **L1/L2 Regularization**: Adds penalty terms to the loss to discourage large weights.
- **Early stopping**: Stops training once validation accuracy stops improving.

6. Optimization Algorithms:

- Optimizers adjust model weights based on gradients of the loss function.
- **SGD (Stochastic Gradient Descent)**: Simple and efficient for large datasets.
- **Adam (Adaptive Moment Estimation)**: Combines momentum and adaptive learning rate.
- The choice of optimizer impacts convergence speed and final accuracy.

7. Batch Normalization:

- Normalizes outputs of layers to reduce internal covariate shift.
- Helps in **faster training, higher learning rates, and better generalization**.

8. Modular and Reusable Design:

- Models are built as reusable components or layers (e.g., in frameworks like TensorFlow or PyTorch).
- Encourages rapid experimentation and testing.

Following these architectural principles helps in building scalable, efficient, and accurate deep learning models suitable for a wide range of applications.

Q4. What is TensorFlow? Explain its features and significance in deep learning.

TensorFlow is an open-source library developed by **Google Brain Team** for performing **deep learning** and **machine learning** tasks. It provides a comprehensive ecosystem of tools for building and deploying models in research and production.

1. Open-Source and Multi-Platform Support:

- TensorFlow is open-source and free to use.
- It supports a wide range of platforms including **Windows, Linux, macOS**, and mobile platforms like **Android** and **iOS**.

- This flexibility makes it suitable for both academic and commercial use.

2. Computational Graph Architecture:

- TensorFlow uses **data flow graphs** to represent computation.
- Each **node** in the graph represents an operation (e.g., addition, multiplication), and **edges** represent tensors (multi-dimensional arrays).
- This approach allows efficient optimization and parallel computation across devices.

3. Eager Execution:

- Eager execution is a flexible mode where operations are executed immediately.
- It allows for easier debugging, step-by-step evaluation, and dynamic model building.
- Ideal for beginners and research-based applications.

4. Keras – High-Level API:

- TensorFlow integrates **Keras**, a simple and intuitive API to build neural networks quickly.
- Keras allows for fast experimentation by reducing code complexity.
- It supports commonly used layers, optimizers, losses, and metrics out of the box.

5. TensorBoard – Visualization Tool:

- TensorBoard is used to **visualize** model training in real-time.
- It helps monitor loss and accuracy, view the computation graph, and debug gradients or parameters.
- This improves model interpretability and optimization.

6. Deployment Tools – Lite, JS, Serving: TensorFlow provides tools for model deployment on various platforms:

- **TensorFlow Lite** for mobile and embedded devices.
- **TensorFlow.js** for running models in the browser using JavaScript.
- **TensorFlow Serving** for deploying models in production environments (e.g., web services).

7. GPU/TPU Support and Scalability:

- TensorFlow supports parallel training on **GPUs** and **TPUs (Tensor Processing Units)**.
- It also allows **distributed training** across multiple machines, making it highly scalable for large datasets and deep architectures.

8. Strong Ecosystem and Community: TensorFlow has a rich ecosystem with additional libraries like:

- **TF Datasets** for data loading.
- **TF Hub** for reusable model components.
- **TF Model Garden** for pre-trained models.
- Backed by Google and an active community, it offers strong support and frequent updates.

Significance in Deep Learning:

- TensorFlow simplifies deep learning development through high-level APIs and modular design.
- It supports **end-to-end workflows** from data preprocessing, model design, training, evaluation, and deployment.
- Used in a variety of fields such as: **Healthcare** (e.g., cancer detection), **Finance** (e.g., fraud detection), **NLP** (e.g., chatbots, machine translation), **Computer Vision** (e.g., facial recognition, autonomous driving).

TensorFlow is a powerful, flexible, and production-ready deep learning framework. Its wide range of features, scalability, and support for model deployment make it an essential tool for modern AI development.

Q5. Describe the role of deep learning in time series, image, and video data analysis.

Deep learning plays a significant role in analyzing different types of data such as time series, images, and videos, enabling accurate predictions, classification, and feature extraction that traditional methods often struggle with.

1. Time Series Data Analysis:

- Time series data is sequential and temporal in nature (e.g., stock prices, weather patterns).
- Deep learning models like **Recurrent Neural Networks (RNNs)** and **Long Short-Term Memory (LSTM)** networks are well-suited because they can capture temporal dependencies and patterns over time.
- Applications include **forecasting future values**, **anomaly detection**, and **speech recognition**.
- Example: Predicting electricity demand based on past consumption data.

2. Image Data Analysis:

- Images are represented as pixel grids with spatial information.
- **Convolutional Neural Networks (CNNs)** are specialized deep learning models that automatically extract hierarchical features such as edges, shapes, and objects.
- CNNs excel in tasks like **image classification**, **object detection**, and **image segmentation**.
- Example: Medical imaging to detect tumors or abnormalities.

3. Video Data Analysis:

- Videos are sequences of images with temporal and spatial information.
- Deep learning combines CNNs (for spatial features) and RNNs or 3D CNNs (for temporal dynamics) to analyze videos.
- Applications include **action recognition**, **video summarization**, and **gesture detection**.
- Example: Surveillance systems detecting suspicious activities from video footage.

Deep learning models provide powerful tools for processing complex data types by automatically learning features from raw data. Their ability to model spatial and temporal dependencies has led to breakthroughs in time series forecasting, image recognition, and video understanding.

Q6. What is Deep Learning? What was the role of deep learning in AlphaGo's success?

Deep Learning is a subset of machine learning that uses multi-layered artificial neural networks to model and understand complex patterns in data. It enables computers to learn hierarchical representations, automatically extracting features from raw input through several layers of abstraction.

Deep Learning:

- Uses **deep neural networks** with many hidden layers.
- Capable of learning **complex non-linear relationships**.
- Can process various data types such as images, text, speech, and time series.
- Employs techniques like **backpropagation** for training networks.
- Has led to significant improvements in tasks like image recognition, natural language processing, and reinforcement learning.

Role of Deep Learning in AlphaGo's Success:

- **AlphaGo** is an AI program developed by DeepMind to play the complex board game Go.
- Deep learning was critical in two parts of AlphaGo's architecture:
 - **Policy Network:** Deep neural networks trained to predict the next move by learning from expert human games and self-play.
 - **Value Network:** Estimates the probability of winning from a given game state.
- The networks enabled AlphaGo to evaluate an enormous number of possible board positions efficiently.
- Combined with **Monte Carlo Tree Search (MCTS)**, deep learning helped AlphaGo outperform traditional Go programs and defeated world champions.
- This success demonstrated deep learning's power in mastering tasks previously thought too complex for AI.

Deep learning was central to AlphaGo's ability to learn strategies and evaluate complex board states, marking a major milestone in AI and showcasing the effectiveness of deep neural networks in reinforcement learning tasks.

Q7. List any five real-world applications of Deep Learning.

Deep learning has found applications across many domains due to its ability to learn complex patterns from large datasets. Here are five significant real-world applications:

1. Image and Video Recognition:

- Used in facial recognition systems (e.g., unlocking phones).
- Applied in medical imaging for detecting diseases such as cancer or diabetic retinopathy.
- Used in autonomous vehicles for object detection and scene understanding.

2. Natural Language Processing (NLP):

- Powers virtual assistants like Siri, Alexa, and Google Assistant.
- Enables machine translation services such as Google Translate.
- Supports chatbots and sentiment analysis in customer service.

3. Speech Recognition:

- Converts spoken language into text (e.g., dictation software).
- Used in voice-controlled devices and transcription services.
- Helps in real-time language translation.

4. Recommendation Systems:

- Powers recommendations on platforms like Netflix, Amazon, and YouTube.
- Analyzes user behavior and preferences to suggest relevant content or products.

5. Autonomous Vehicles:

- Enables self-driving cars to perceive the environment using cameras and sensors.
- Processes complex data for navigation, obstacle detection, and decision making.

These applications demonstrate the versatility and impact of deep learning in transforming technology and everyday life.

Q8. Compare and contrast machine learning and deep learning.

Machine learning and deep learning are both subsets of artificial intelligence, but they differ significantly in their approaches, capabilities, and applications.

Aspect	Machine Learning (ML)	Deep Learning (DL)
Definition	A method of teaching computers to learn from data using algorithms like decision trees, SVM, etc.	A specialized subset of ML that uses deep neural networks with many layers to learn from large datasets.
Feature Engineering	Requires manual feature extraction and selection by domain experts.	Automatically extracts features from raw data using multiple layers.
Data Requirements	Works well with smaller to medium-sized datasets.	Requires large amounts of data to perform effectively.
Computational Power	Less computationally intensive; can run on CPUs.	Requires high computational power, often GPUs or TPUs.
Performance	Effective for simple to moderately complex problems.	Excels at solving complex problems like image and speech recognition.
Interpretability	More interpretable and transparent models (e.g., decision trees).	Often considered a “black box” due to complex internal workings.
Applications	Spam detection, credit scoring, customer churn prediction.	Image recognition, natural language processing, autonomous driving.

While machine learning is suitable for many traditional data problems, deep learning provides powerful tools for handling unstructured data and complex tasks but requires more data and resources.

UNIT II: Fundamentals of Neural Network

Q1. Describe the structure and working of an artificial neuron with a diagram.

An artificial neuron is a fundamental unit of a neural network inspired by the biological neuron in the human brain. It mimics the process of receiving inputs, processing them, and generating an output signal based on certain conditions.

Structure:


- **Inputs ($x_1, x_2, \dots, x_n$):** These are numerical signals or feature values fed into the neuron.
- **Weights ($w_1, w_2, \dots, w_n$):** Each input is multiplied by a corresponding weight representing its importance.
- **Bias (b):** A constant added to the weighted sum to allow shifting of the activation function.
- **Summation function:** Computes the weighted sum of inputs plus bias:
$$z = \sum_{i=1}^n w_i x_i + b$$
- **Activation function (f):** Applies non-linear transformation to the sum, producing neuron's output: $y = f(z)$

Working:

1. Each input is multiplied by its weight.
2. All weighted inputs are summed along with the bias.
3. The sum is passed through an activation function to introduce non-linearity.
4. The output y is forwarded to the next layer or used as the final result.

Diagram explanation:

- Multiple inputs $x_1, x_2, \dots, x_n$ enter the neuron.
- Each input is multiplied by a weight $w_1, w_2, \dots, w_n$.
- The weighted inputs are summed with bias b inside the summation node.
- The sum is passed to an activation function f .
- The activated output y is generated.



Q2. Explain different learning rules used in neural networks.

Learning in neural networks involves adjusting weights and biases to minimize the error between the predicted output and the actual output. Learning rules define how these adjustments are made based on the input data and error signals. Different learning rules are applied depending on the type of neural network and learning approach.

Types of Learning Rules:

1. **Hebbian Learning Rule:** Based on the principle proposed by Donald Hebb (1949), often summarized as "neurons that fire together, wire together." It is an unsupervised learning rule where the connection strength between two neurons is increased if they are activated simultaneously.
 - The weight update rule can be expressed as:
$$\Delta w_{ij} = \eta \cdot x_i \cdot y_j$$
 - where η is the learning rate, x_i is the input, and y_j is the output of the neuron.
 - This rule is used mainly in associative memory networks and self-organizing maps.
2. **Perceptron Learning Rule:** A supervised learning rule used in single-layer perceptrons. Weights are updated based on the difference between the actual output and the desired output.
 - The update formula is:
$$w_i^{new} = w_i^{old} + \eta(d - y)x_i$$
 - where d is the desired output, y is the actual output, and η is the learning rate.
 - The perceptron learning rule guarantees convergence if the data is linearly separable.

3. **Delta Rule (Widrow-Hoff Rule):** An extension of the perceptron learning rule for networks with continuous activation functions like sigmoid. It uses gradient descent to minimize the mean squared error between actual and desired outputs.
 - The weight update is: $\Delta w_i = \eta(d - y)x_i$
 - This rule forms the basis for more advanced algorithms like backpropagation.
4. **Backpropagation Algorithm:** This is the most widely used learning algorithm for multi-layer neural networks. It's a form of gradient descent that works in two phases:
 - **Forward Pass:** The network processes input data and produces an output.
 - **Backward Pass:** The error (or loss) is calculated at the output layer. The algorithm then propagates this error backward through the network using the chain rule of calculus to calculate the gradient of the loss function with respect to each weight. These gradients indicate the direction and magnitude of the change needed for each weight to reduce the error. The weights and biases are then updated using these gradients. This process is repeated for many iterations until the network's performance is satisfactory
5. **Competitive Learning Rule:** An unsupervised learning rule where neurons compete to become active. Only the neuron with the highest activation (winner) gets to update its weights to match the input pattern. Used in networks like Kohonen's Self-Organizing Maps (SOM).

Q3. What is an activation function? Give examples. Explain its use in neural networks.

An **activation function** is a mathematical function that introduces non-linearity into a neural network. It's a critical component of an artificial neuron that determines the output of the neuron based on the weighted sum of its inputs and a bias

An **activation function** ϕ transforms the input $\text{net} = \sum w_i x_i + b$ into an output $y = \phi(\text{net})$. Without activation functions, the neural network would behave like a linear regression model regardless of the number of layers.

Types and Examples of Activation Functions:

1. **Step Function (Threshold Function):** $\phi(\text{net}) = \begin{cases} 1 & \text{if } \text{net} \geq \theta \\ 0 & \text{if } \text{net} < \theta \end{cases}$ where θ is the threshold. Used in early perceptrons; produces binary output.
2. **Linear Activation Function:** $\phi(\text{net}) = \text{net}$ Output is proportional to input. Limited use because it does not introduce non-linearity.
3. **Sigmoid Function:** $\phi(\text{net}) = \frac{1}{1 + e^{-\text{net}}}$ Outputs values between 0 and 1. Smooth and differentiable, suitable for binary classification.
4. **Hyperbolic Tangent (tanh):** $\phi(\text{net}) = \tanh(\text{net}) = \frac{e^{\text{net}} - e^{-\text{net}}}{e^{\text{net}} + e^{-\text{net}}}$ Outputs values between -1 and 1. Zero-centered, which often helps training.
5. **Rectified Linear Unit (ReLU):** $\phi(\text{net}) = \max(0, \text{net})$ Allows faster and more effective training by introducing sparsity and avoiding vanishing gradients.

Use of Activation Functions in Neural Networks:

- **Non-linearity:** Activation functions enable neural networks to model non-linear relationships, which is essential for tasks like image recognition, natural language processing, etc.
- **Decision Making:** They help decide whether a neuron should be activated or not, based on the input signals.
- **Bounded Output:** Functions like sigmoid or tanh squash the output to a bounded range, making it easier to interpret as probabilities.
- **Gradient-Based Learning:** Differentiable activation functions (sigmoid, tanh, ReLU) enable the use of gradient descent and backpropagation for training deep networks.

Activation functions play a critical role by introducing non-linearity into the neural network, allowing it to learn complex mappings between inputs and outputs. Different activation functions serve various purposes, and the choice depends on the network architecture and task.

Q4. What is a perceptron? Describe perceptron learning algorithm.

A **perceptron** is the simplest form of an artificial neuron and a type of linear classifier. It was developed in the 1950s and is capable of solving problems that are **linearly separable**. It takes multiple inputs, multiplies them by weights, sums them up, and passes the result through a step-function activation function to produce a binary output (typically 0 or 1).

What is a Perceptron?

- It consists of one layer of artificial neurons (often just one neuron).
- Inputs are multiplied by weights, summed, and passed through a step activation function (threshold function) to produce a binary output (0 or 1).
- It can classify linearly separable data by finding a hyperplane separating the classes.

Mathematical Model:
$$y = \phi \left(\sum_{i=1}^n w_i x_i + b \right)$$

where ϕ is the step function:
$$\phi(\text{net}) = \begin{cases} 1 & \text{if } \text{net} \geq 0 \\ 0 & \text{if } \text{net} < 0 \end{cases}$$

Perceptron Learning Algorithm: The perceptron learning algorithm iteratively adjusts weights to minimize classification errors on training data.

Steps:

1. **Initialize weights and bias:** Set weights w_i and bias b to small random values or zeros.
2. **Input training sample:** Present an input vector $x = (x_1, x_2, \dots, x_n)$ with desired output $d \in \{0,1\}$.
3. **Compute output:** Calculate weighted sum $\text{net} = \sum w_i x_i + b$, and apply the step activation function to get predicted output y .
4. **Update weights:** If output y differs from desired output d , update weights and bias using:

where η is the learning rate (a small positive constant).

$$\begin{aligned} w_i &\leftarrow w_i + \eta(d - y)x_i \\ b &\leftarrow b + \eta(d - y) \end{aligned}$$

5. **Repeat for all training samples:** Iterate over the entire training set multiple times until convergence or for a fixed number of epochs.

Working Principle:

- If the perceptron predicts correctly, weights remain unchanged.
- If it misclassifies, weights are adjusted to reduce error, nudging the decision boundary in the right direction.
- The algorithm converges if the data is linearly separable.

Limitations: Can only solve linearly separable problems like AND/OR. Cannot solve non-linear problems like XOR.

The perceptron is a simple neural model for binary classification. The perceptron learning algorithm iteratively updates weights to correctly classify input patterns by minimizing error using a step activation function.

Q5. Compare linear and non-linear activation functions with suitable examples.

Activation functions determine the output of a neuron given its input. They can be broadly classified into **linear** and **non-linear** types, each with distinct properties and uses.

1. Linear Activation Function: A linear activation function outputs a value proportional to its input. Mathematically: $\phi(\text{net}) = k \cdot \text{net}$ where k is a constant, often set to 1.

- **Characteristics:** Output is unbounded and directly proportional to input. Simple, computationally efficient.
- **Limitations:** No matter how many layers the network has, stacking linear activation functions results in a model that behaves like a single-layer linear model.
 - Cannot model complex non-linear relationships.
 - Makes the network equivalent to a linear regression.
- **Use Case:** Primarily used in the output layer for regression tasks where continuous outputs are needed.

2. Non-linear Activation Function: Non-linear activation functions produce outputs that are non-linear transformations of their inputs. This enables the neural network to learn complex, non-linear mappings.

- **Sigmoid:** $\phi(net) = \frac{1}{1 + e^{-net}}$ Output range: (0, 1)
- **Hyperbolic Tangent (tanh):** $\phi(net) = \tanh(net) = \frac{e^{net} - e^{-net}}{e^{net} + e^{-net}}$ Output range: (-1, 1)
- **Rectified Linear Unit (ReLU):** $\phi(net) = \max(0, net)$ Output range: [0, ∞)
- **Characteristics:** Introduces non-linearity to the model.
 - Allows the network to learn and represent complex functions.
 - Differentiable (except at some points like ReLU at zero), enabling gradient-based optimization.
 - Helps in solving real-world problems with complex decision boundaries.
- **Use Case:** Used in hidden layers of neural networks for classification, regression, and deep learning tasks.

Aspect	Linear Activation	Non-linear Activation
Output	Proportional to input	Non-linear transformation of input
Complexity	Simple	Complex
Network Capability	Only linear problems	Can model complex, non-linear problems
Stacking Effect	Equivalent to single layer	Enables deep networks to learn features
Differentiability	Yes	Yes (mostly)
Examples	$y = x$	Sigmoid, tanh, ReLU
Usage	Output layer for regression	Hidden layers and classification output

Linear activation functions are simple but limited to modeling linear relationships, whereas non-linear activation functions enable neural networks to learn complex patterns and solve real-world problems effectively. Hence, non-linear activations are preferred in most hidden layers.

Q6. Differentiate between single-layer and multi-layer feed-forward networks.

Feed-forward neural networks are the simplest type of artificial neural networks where information flows only in one direction—from input to output—without loops. Depending on the number of layers, they are classified as single-layer or multi-layer networks.

Single-Layer Feed-Forward Network:

- **Structure:** Consists of only one layer of neurons (output layer) directly connected to the input layer. There are no hidden layers.
- **Operation:** The inputs are directly connected to the output neurons, which compute the weighted sum and apply an activation function (usually a step or linear function).
- **Learning Capability:** Can only solve **linearly separable problems** because it implements a linear decision boundary.
- **Examples:** Perceptron, Single-layer Adaline (Adaptive Linear Neuron)
- **Limitations:** Cannot solve complex problems like XOR which are not linearly separable.

Multi-Layer Feed-Forward Network (MLP - Multi-Layer Perceptron):

- **Structure:** Consists of an input layer, one or more hidden layers, and an output layer. Each layer consists of multiple neurons.
- **Operation:** Information flows from input to output passing through hidden layers, where each neuron performs a weighted sum followed by a non-linear activation function.

- **Learning Capability:** Can solve **both linear and non-linear problems** by learning complex mappings between inputs and outputs.
- **Examples:** Multi-layer perceptron trained using backpropagation algorithm.
- **Advantages:** Can approximate any continuous function (Universal Approximation Theorem). Suitable for tasks like image recognition, speech processing, and other complex classifications.
- **Disadvantages:** More computationally intensive due to multiple layers and parameters.

Feature	Single-Layer Feed-Forward Network	Multi-Layer Feed-Forward Network
Number of Layers	One (Output layer only)	Multiple (One or more hidden layers + output)
Complexity	Simple	Complex
Learning Capability	Can solve only linearly separable problems	Can solve both linear and non-linear problems
Activation Function	Usually linear or step function	Non-linear functions like sigmoid, ReLU
Training Algorithm	Perceptron learning rule	Backpropagation algorithm
Example Applications	Basic classification	Image processing, NLP, function approximation
Computational Cost	Low	High

Q7. Explain the working of a Recurrent Neural Network (RNN) with its applications.

Recurrent Neural Networks (RNNs) are a class of artificial neural networks designed to handle sequential or time-series data. Unlike feed-forward networks, RNNs have loops that allow information to persist, making them suitable for tasks where context and order matter.

Structure of RNN: Consists of input layer, one or more recurrent hidden layers, and output layer. Each neuron in the hidden layer receives input from the current time step and from its own output at the previous time step (feedback loop).

Working Principle:

- At time step t , the RNN takes input vector x_t and the hidden state from the previous time step h_{t-1} .
- The hidden state is updated as: $h_t = \phi(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$ where W_{xh} and W_{hh} are weight matrices, b_h is bias, and ϕ is the activation function (typically tanh or ReLU).
- The output at time t is computed as: $y_t = \psi(W_{hy}h_t + b_y)$ where W_{hy} is output weight matrix and ψ is usually softmax or linear function.
- This recurrence allows the network to maintain a **memory** of previous inputs via the hidden state h_t , enabling it to process sequences.

Key Characteristics:

- **Temporal Dynamics:** RNNs capture temporal dependencies in sequential data.
- **Parameter Sharing:** The same weights are used at every time step, reducing the number of parameters.
- **Backpropagation Through Time (BPTT):** A variant of backpropagation is used to train RNNs by unfolding the network across time steps.

Applications of RNNs:

1. **Natural Language Processing (NLP):** Language modeling, text generation, machine translation, and sentiment analysis.
2. **Speech Recognition:** Understanding spoken language by modeling audio sequences.
3. **Time Series Prediction:** Stock price forecasting, weather prediction, and sensor data analysis.
4. **Music Generation:** Composing music by learning patterns in note sequences.
5. **Video Analysis:** Activity recognition and video captioning by analyzing frames over time.

Limitations:

- **Vanishing and Exploding Gradients:** Difficulties in learning long-term dependencies.
- Addressed by advanced RNN variants like LSTM (Long Short-Term Memory) and GRU (Gated Recurrent Unit).

Recurrent Neural Networks are powerful models for sequential data that use feedback connections to remember past information. Their ability to model time-dependent patterns makes them invaluable in many real-world applications involving sequences.

Q8. Discuss various types of neural network architectures in detail.

Neural networks come in various architectures designed to process different types of data such as static, sequential, or spatial data. Each architecture has unique features and is suited to specific problem types.

1. **Feed-Forward Neural Networks (FFNN):** Information flows in one direction from input to output without cycles. Comprises input layer, one or more hidden layers, and an output layer. Suitable for static data classification and regression problems. Examples include single-layer perceptron and multi-layer perceptron (MLP).
2. **Convolutional Neural Networks (CNN):** Specialized for processing grid-like data such as images. Uses convolutional layers to automatically learn spatial features like edges, textures, and objects. Includes convolutional, pooling, and fully connected layers. Widely used in image recognition, video analysis, and computer vision tasks.
3. **Recurrent Neural Networks (RNN):** Designed for sequential data processing by maintaining a hidden state that captures information from previous time steps. Effective for tasks where temporal context is important. Traditional RNNs suffer from vanishing and exploding gradients which limit long-term memory.
4. **Long Short-Term Memory (LSTM) Networks:** An advanced variant of RNN designed to overcome limitations of standard RNNs. Contains specialized memory cells and gates (forget, input, output) that control information flow. Capable of learning long-term dependencies in sequences. Used in language modeling, speech recognition, time series forecasting.
5. **Gated Recurrent Unit (GRU) Networks:** Similar to LSTM but with a simpler structure using only two gates: reset and update gates. Requires fewer parameters and is faster to train than LSTM while still managing long-term dependencies well. Applied in similar domains as LSTM.
6. **Radial Basis Function Networks (RBFN):** Employ radial basis functions as activation in hidden layers. Typically three-layered (input, hidden with RBF neurons, output). Used for function approximation, regression, and classification.
7. **Modular Neural Networks (MNN):** Composed of multiple independent modules or subnetworks, each solving part of a problem. Modules' outputs are combined for final decision-making. Useful for complex problems decomposed into simpler tasks.
8. **Self-Organizing Maps (SOM):** Unsupervised networks that map high-dimensional input data into low-dimensional (usually 2D) maps preserving topological properties. Used for clustering, visualization, and dimensionality reduction.
9. **Autoencoders:** Unsupervised neural networks designed to learn compressed representations of data. Consist of encoder and decoder parts. Useful for dimensionality reduction, anomaly detection, and denoising.

Neural network architectures vary widely, from simple feed-forward networks for static data to complex recurrent networks like LSTM and GRU for sequential data. CNNs excel in spatial data analysis, while unsupervised architectures like SOM and autoencoders serve clustering and compression roles. Selecting the right architecture depends on the data type and task requirements.

Q9. Describe how neural networks are used in music genre classification.

Music genre classification is the task of categorizing music tracks into predefined genres (e.g., rock, jazz, classical) based on audio characteristics. Neural networks are widely used for this purpose due to their ability to learn complex patterns from raw or processed audio data.

Process of Music Genre Classification Using Neural Networks:

1. **Feature Extraction:** Raw audio signals are transformed into feature representations such as Mel-Frequency Cepstral Coefficients (MFCCs), chroma features, spectral contrast, and rhythm patterns. These features capture timbre, pitch, rhythm, and other characteristics essential for genre discrimination.

2. **Input to Neural Network:** Extracted features are used as inputs to the neural network. Depending on architecture, inputs can be raw audio frames (in CNNs) or sequences of feature vectors (in RNNs).
3. **Neural Network Architecture:**
 - **Feed-Forward Networks (MLP):** Can be used with handcrafted features for classification.
 - **Convolutional Neural Networks (CNN):** Process spectrograms or other 2D representations to capture local and hierarchical audio patterns.
 - **Recurrent Neural Networks (RNN) / LSTM:** Model temporal dependencies and sequence information in music, capturing rhythm and progression.
4. **Training the Network:** The network is trained using labeled datasets where each music sample is tagged with a genre. Loss functions (e.g., cross-entropy) are minimized using gradient descent and backpropagation.
5. **Classification:** After training, the neural network predicts genre of unseen music samples on learned patterns.

Advantages of Neural Networks for Music Genre Classification:

- Ability to learn complex and subtle audio features automatically.
- Effective in modeling both spectral and temporal aspects of music.
- Superior performance compared to traditional machine learning methods relying on manual feature engineering.

Applications:

- Music streaming services for personalized recommendations.
- Automatic organization and tagging of large music libraries.
- Music information retrieval and analysis.

Neural networks are powerful tools for music genre classification. By extracting meaningful audio features and leveraging architectures like CNNs and RNNs, they learn to discriminate between genres accurately, enabling various practical applications in music technology.

Q10. Explain the role of activation functions in deep learning networks with diagrams.

Activation functions introduce non-linearity into neural networks, which is critical for deep learning models to learn complex mappings between inputs and outputs. Without activation functions, the network behaves like a linear model, regardless of its depth.

Role of Activation Functions in Deep Learning:


1. **Introducing Non-Linearity:** Real-world data often involve non-linear relationships. Activation functions enable neurons to learn these by applying non-linear transformations to inputs. This allows deep networks with multiple layers to approximate any continuous function (Universal Approximation Theorem).
2. **Controlling Signal Flow:** Activation functions determine whether a neuron should be activated based on its input, allowing selective passing of information.
3. **Facilitating Gradient-Based Optimization:** Activation functions must be differentiable so that algorithms like backpropagation can compute gradients and update weights effectively.
4. **Preventing Output Saturation:** Functions like ReLU reduce the problem of vanishing gradients common with sigmoid/tanh in deep networks, enabling faster and more stable training.

Common Activation Functions Used in Deep Learning:


Activation Function	Formula	Output Range	Characteristics
Sigmoid	$\sigma(x) = 1 / (1 + e^{-x})$	(0, 1)	Smooth, saturates at extremes
Tanh	$\tanh(x) = (e^x - e^{-x}) / (e^x + e^{-x})$	(-1, 1)	Zero-centered, saturates at extremes

ReLU	$f(x) = \max(0, x)$	$[0, \infty)$	Sparse activation, efficient
Leaky ReLU	$f(x) = \max(\alpha x, x)$ (α small, e.g., 0.01)	$(-\infty, \infty)$	Allows small gradient when inactive

Sigmoid Activation Function



ReLU Activation Function



Q11. Write a short note on how neural networks learn patterns from data.

Neural networks learn by adjusting their internal parameters (weights and biases) to minimize the difference between predicted outputs and actual targets in the training data. This learning process enables the network to recognize underlying patterns and make accurate predictions on new, unseen data.

How Neural Networks Learn Patterns:

- Initialization:** Weights and biases are initialized, usually with small random values, before training begins.
- Forward Propagation:**
 - Input data is passed through the network layer by layer.
 - Each neuron calculates a weighted sum of its inputs and applies an activation function to produce output.
 - The final output layer produces the network's prediction.
- Error Calculation:**
 - The predicted output is compared with the actual target using a loss function (e.g., mean squared error for regression, cross-entropy for classification).
 - The loss quantifies how well the network predicts the data.
- Backward Propagation (Backpropagation):**
 - The error is propagated backward through the network.
 - Gradients of the loss function with respect to each weight are computed using calculus (chain rule).
 - This determines how much each weight contributes to the error.
- Weight Update:**
 - Using optimization algorithms like gradient descent, weights are updated to reduce the error.
 - Updates are typically proportional to the negative gradient scaled by a learning rate.
- Iteration:** Steps 2 to 5 are repeated over many epochs (complete passes through the training data) until the network converges or reaches a satisfactory accuracy.

Learning Outcome: Through this iterative process, the network tunes its weights to capture patterns and relationships in the input data, allowing it to generalize and make predictions on new inputs.

UNIT III: Convolutional Neural Network (CNN)

1. What is the motivation behind Convolutional Neural Networks (CNNs)?

The primary motivation behind Convolutional Neural Networks (CNNs) is to efficiently process data with a grid-like topology, such as images, by overcoming the limitations of traditional feedforward neural networks (like Multi-Layer Perceptrons, or MLPs). MLPs are not well-suited for image processing for two main reasons:

- High Dimensionality:** A simple 256x256 pixel image has over 65,000 pixels. If this is fed into an MLP, the first hidden layer would have an enormous number of weights, leading to a massive number of parameters. This is computationally expensive and makes the network highly prone to **overfitting**.
- Loss of Spatial Structure:** MLPs require the input image to be flattened into a 1D vector. This process completely discards the spatial information of the image (i.e., which pixels are near each other). However, in images, nearby pixels are highly correlated and form meaningful local patterns like edges, corners, and textures.

CNNs were motivated by the need to address these issues. They are inspired by the human visual cortex and are designed to:

- Preserve Spatial Hierarchy:** process images in their 2D form, preserving spatial relationships between pixels.
- Learn Hierarchical Features:** CNNs automatically learn a hierarchy of features. Early layers learn simple features like edges and colors, while deeper layers combine these to learn more complex features like shapes, textures, and eventually, objects.
- Be Computationally Efficient:** They use two key principles—**local receptive fields** and **parameter sharing**—to drastically reduce the number of parameters compared to a fully connected network, making them more efficient and less prone to overfitting.

2. Differentiate between convolutional layer and pooling layer in CNN.

Feature	Convolutional Layer	Pooling Layer
Primary Purpose	Feature Extraction. To detect patterns like edges, corners, and textures in the input feature map.	Downsampling. To reduce the spatial dimensions (width and height) of the feature map.
Operation	Performs a convolution operation, which is a linear operation involving a dot product between the filter and local regions of the input.	Performs a non-linear aggregation operation, typically Max Pooling (taking the maximum value) or Average Pooling .
Parameters	Has learnable parameters: the weights and biases of its filters (kernels). These are updated during training via backpropagation.	Has no learnable parameters. The operation (e.g., finding the max in a 2x2 window) is fixed.
Output Size	The output feature map's size is controlled by filter size, stride, and padding. It can be the same size, smaller, or even larger than the input.	The output feature map is always smaller in spatial dimensions than the input.

Key Contribution	It identifies <i>what</i> features are present in the input and <i>where</i> they are located, creating feature maps.	It makes the network more computationally efficient and provides a degree of translation invariance , making the feature detection more robust to small shifts in the input.
-------------------------	---	---

3. Explain the concept of parameter sharing in CNNs with an example.

Parameter sharing is a core concept in CNNs that dramatically reduces the number of parameters in the model. The key idea is that a feature detector (a filter) that is useful in one part of the image is likely also useful in another part of the image.

Concept: Instead of learning a separate set of weights for every single location in the input image, a CNN learns a single set of weights (a filter) and applies it across the entire image. This means the same filter is used to detect a specific feature (e.g., a vertical edge) regardless of where it appears. All the neurons in a given feature map share the same set of weights and bias.

Example: Imagine we want to detect vertical edges in a 5x5 grayscale image. We can design a 3x3 filter (or kernel) for this task: Vertical Edge Filter=111000-1-1-1

1. **Without Parameter Sharing (Hypothetical):** In a fully connected network, to process a 3x3 patch, we would need 9 unique weights. To process the entire 5x5 image by looking at every 3x3 patch, we would need a different set of 9 weights for each patch, resulting in a very large number of parameters.
2. **With Parameter Sharing (CNN approach):** The CNN uses this *single* 3x3 filter and "slides" it over the entire 5x5 input image. The *same* 9 weights of the filter are used at every position. The filter performs a convolution operation at each location to produce an output value in the feature map.

Benefits:

- **Massive Parameter Reduction:** We only need to learn the 9 weights of the filter (plus one bias term), instead of thousands. This makes the network much more efficient to train and less likely to overfit.
- **Translation Invariance:** The network can detect a feature (like a vertical edge) no matter where it is in the image because the same detector is used everywhere.

4. What are filters (kernels) in CNN? How do they help in feature extraction?

In a CNN, **filters** (also known as **kernels**) are small, multi-dimensional arrays of learnable weights. They are the core components responsible for detecting patterns and features in the input data.

A filter is essentially a feature detector. During the convolution operation, the filter slides over the input image (or the feature map from a previous layer), and at each location, it computes a dot product between its own weights and the input pixels it is currently on.

How they help in feature extraction:

1. **Specialized Detectors:** Each filter in a convolutional layer learns to recognize a specific low-level feature. For example, in the initial layers of a network trained on images, different filters might learn to detect:
 - Horizontal, vertical, or diagonal edges
 - Specific colors or color gradients
 - Simple textures like grids or dots
2. **Creating Feature Maps:** The output of applying a filter across an entire input is called a **feature map** or an

activation map. This map highlights the locations in the input where the filter's specific feature was detected. A high activation value in the feature map means that the feature the filter is looking for is strongly present at that location in the input.

3. **Hierarchical Feature Learning:** A CNN uses multiple filters in each layer, creating a stack of feature maps. These maps are then used as the input to the next layer. The filters in deeper layers learn to combine the simple features detected by earlier layers into more complex patterns. For example, filters in a deeper layer might learn to detect:
- Corners (by combining horizontal and vertical edges)
 - Shapes (by combining corners and edges)
 - Parts of objects (like an eye or a car wheel)

Through this process, the network builds a rich, hierarchical representation of the input, moving from simple pixels to abstract concepts.

5. Compare fully connected layers and convolutional layers.

Feature	Fully Connected Layer (FCL)	Convolutional Layer (Conv Layer)
Neuron Connectivity	Dense. Every neuron in the layer is connected to <i>every</i> neuron in the previous layer.	Local. Neurons are only connected to a small, local region of the input (the receptive field).
Input Shape	Assumes the input is a 1D vector . Any multi-dimensional input must be flattened first.	Preserves the input's spatial structure (e.g., 2D for images, 3D for videos).
Parameter Sharing	No parameter sharing. Each connection has its own unique weight, leading to a very large number of parameters.	Extensive parameter sharing. The same filter (set of weights) is used across the entire input, drastically reducing the number of parameters.
Primary Use Case	Typically used at the end of a CNN for classification or regression, after features have been extracted.	Used in the initial and middle stages of a CNN for hierarchical feature extraction.
Strengths	Can learn global, non-linear combinations of high-level features provided by preceding layers.	Computationally efficient, less prone to overfitting on spatial data, and excellent at learning spatial hierarchies of features.
Weaknesses	Computationally expensive, prone to overfitting, and discards all spatial information from the input.	Not well-suited for non-spatial data or for learning global patterns in the final classification stage.

6. Explain the role of dropout regularization in CNNs.

Dropout is a powerful regularization technique used to prevent overfitting in neural networks, including CNNs.

Concept: During the training phase, dropout randomly sets the output of a certain percentage of neurons in a layer to zero for each forward pass. The neurons to be "dropped" are chosen randomly for each training batch. This means

that for each training step, the network is trained with a different, thinned-out architecture.

Role and Mechanism:

1. **Prevents Co-adaptation:** By randomly deactivating neurons, dropout ensures that no single neuron becomes overly specialized or reliant on the presence of another specific neuron. Neurons must learn more robust features that are useful in conjunction with different random subsets of other neurons. This breaks up complex co-adaptations where the network might be memorizing noise in the training data.
2. **Encourages Redundancy:** Since any neuron can be dropped out, the network is forced to learn redundant representations of the data. It cannot rely on one specific pathway to get the right answer; instead, it learns multiple, independent ways to detect features.
3. **Acts as Model Averaging:** Training a network with dropout is analogous to training a large number of different, smaller neural networks simultaneously. At test time, dropout is turned off, and the full network is used. This is equivalent to taking an average of the predictions of all the thinned-out networks, which typically results in better generalization and performance on unseen data.

In CNNs, dropout is most commonly applied to the fully connected layers at the end of the network, as these layers contain the majority of the parameters and are therefore most susceptible to overfitting.

7. What is weight decay (L2 regularization)? How does it prevent overfitting?

Weight decay, also known as **L2 regularization**, is a technique used to prevent overfitting by adding a penalty term to the network's loss function.

What it is: The standard loss function (e.g., Cross-Entropy) aims to minimize the error between the network's predictions and the true labels. L2 regularization adds a new term to this loss function that penalizes large weight values. The modified loss function becomes: $L_{new}(\theta) = L_{original}(\theta) + 2m\lambda w \sum w^2$

Where:

- $L_{original}(\theta)$ is the original loss function.
- $\sum w^2$ is the sum of the squares of all the weights (w) in the network.
- λ (lambda) is the regularization hyperparameter, which controls the strength of the penalty. A higher λ means a stronger penalty.
- m is the number of training examples.

How it prevents overfitting:

1. **Penalizes Model Complexity:** A model that has overfit the training data often has very large weight values. This happens because the model tries to perfectly fit every data point, including the noise, by making its decision boundaries highly complex. By adding a penalty for large weights, L2 regularization discourages the model from learning such complex, high-variance functions.
2. **Encourages Smaller Weights:** During training (backpropagation), the gradient of this new loss function includes a term that pushes the weights towards zero. This "decays" the weights, forcing the network to maintain smaller values.
3. **Creates a "Simpler" Model:** A model with smaller weights is essentially a simpler model. It is less sensitive to small changes in the input features. This results in a "smoother" decision boundary and a model that relies on a broader set of weak features rather than a few highly specialized, strong features. This makes the model generalize better to new, unseen data.

8. Write short notes on max pooling in CNNs.

Max Pooling is a downsampling operation commonly used in convolutional neural networks between successive convolutional layers. Its primary goal is to reduce the spatial dimensions (width and height) of the feature maps, which has several key benefits.

Process:

1. A pooling window (or kernel), typically of size 2x2, is defined along with a stride, usually 2.
2. This window slides over each feature map from the previous layer.
3. For each region covered by the window, the maximum pixel value within that region is taken as the output. All other values in the window are discarded.
4. This creates a new, smaller feature map that retains the most prominent features from the original map.

Example: For a 2x2 max pooling operation with a stride of 2 on a 4x4 feature map:

13908526241567382x2 Max Pool(max(1,8,3,5)max(9,2,0,6)max(2,6,4,7)max(1,3,5,8))=(8978)

Key Roles and Advantages:

- **Dimensionality Reduction:** It significantly reduces the number of parameters and computations required in subsequent layers, making the network faster and more memory-efficient.
- **Translation Invariance:** It makes the feature detection more robust to small shifts and distortions in the input image. If a feature moves slightly within the pooling window, the output of the pooling layer will likely remain the same because the maximum value is preserved.
- **Overfitting Control:** By summarizing features in a neighborhood, it helps to prevent the network from overfitting to the exact spatial locations of features in the training data.

9. Explain the process of backpropagation in CNNs.

Backpropagation is the algorithm used to train a CNN by iteratively adjusting the learnable parameters (weights and biases of the filters) to minimize the loss function. It is an application of the chain rule from calculus to compute the gradient of the loss with respect to each parameter. The process involves two main passes: a forward pass and a backward pass.

1. **Forward Pass:** An input image is fed into the network. It passes sequentially through the layers: convolution, activation (e.g., ReLU), pooling, and fully connected layers. Each layer transforms the input from the previous layer, and the process continues until the final layer produces an output prediction (e.g., a vector of class probabilities).
2. **Calculate Loss:** The network's prediction is compared to the ground-truth label using a loss function (e.g., Cross-Entropy Loss). The loss is a single scalar value that quantifies how wrong the network's prediction was.
3. **Backward Pass (Backpropagation of Error):** The goal is to calculate the gradient of the loss function with respect to every weight and bias in the network ($\partial W \partial L$). This process starts from the final layer and moves backward towards the input layer, applying the chain rule at each step.
 - **For Fully Connected Layers:** The process is the same as in a standard MLP.
 - **For Pooling Layers:** Since pooling (e.g., max pooling) has no parameters, the error is simply passed back to the neuron that contributed the maximum value during the forward pass. The gradients for all other neurons in that patch are zero.
 - **For Convolutional Layers:** This is the most complex step. The gradient of the loss is propagated back to the filter weights. Because the same filter is used across multiple locations (parameter sharing), the gradients from all these locations are summed up to get the final gradient for that filter.
4. **Weight Update:**
 - Once the gradients for all parameters are calculated, an optimization algorithm like Stochastic Gradient Descent (SGD) or Adam is used to update the weights and biases.
 - The update rule is: $W_{\text{new}} = W_{\text{old}} - \eta \partial W \partial L$, where η is the learning rate.
 - The parameters are adjusted in the direction that most effectively reduces the loss.

This entire cycle of forward pass, loss calculation, backward pass, and weight update is repeated for many epochs until the network's performance converges.

10. Write a short note on LeNet architecture and its applications.

LeNet (specifically LeNet-5) is one of the earliest and most influential Convolutional Neural Network architectures, developed by Yann LeCun and his colleagues in the 1990s. It is considered a pioneering work that demonstrated the power of CNNs for document recognition tasks.

Architecture: LeNet-5 has a simple and effective architecture that established the standard template for many modern CNNs: a sequence of convolutional and pooling layers, followed by fully connected layers for classification. The typical LeNet-5 architecture consists of 7 layers (excluding the input):

1. **C1 (Convolutional Layer):** Uses 6 filters of size 5x5 to create 6 feature maps.
2. **S2 (Subsampling/Pooling Layer):** Performs average pooling with a 2x2 window to downsample feature maps.
3. **C3 (Convolutional Layer):** A more complex layer with 16 filters of size 5x5, which combines inputs from the S2 layer's feature maps.
4. **S4 (Subsampling/Pooling Layer):** Another average pooling layer with a 2x2 window.
5. **C5 (Convolutional Layer):** Has 120 filters of size 5x5. This layer is sometimes viewed as a fully connected layer because its input and filter sizes match.
6. **F6 (Fully Connected Layer):** Contains 84 neurons.
7. **Output Layer:** A final fully connected layer with 10 output neurons (one for each digit 0-9), often using a softmax activation for classification.

LeNet used tanh or sigmoid as its activation functions, unlike modern networks that primarily use ReLU.

Applications: LeNet was originally designed and famously used for handwritten digit recognition. Its primary commercial application was in reading zip codes on mail and recognizing numbers on bank checks, tasks where it achieved very high accuracy and demonstrated its real-world viability. Although modern architectures are far more complex, the fundamental principles pioneered by LeNet form the foundation of virtually all modern computer vision applications.

11. What is PixelNet? Mention its applications.

PixelNet is a deep neural network architecture designed for dense, per-pixel prediction tasks, most notably **semantic segmentation**. Unlike other popular architectures for segmentation like Fully Convolutional Networks (FCNs), which use upsampling or deconvolution layers to generate a full-resolution output, PixelNet takes a different approach.

Concept: to make an independent classification decision for each pixel in the input image. To do this, for every pixel, it gathers features from different depths of a pre-trained base CNN (like VGG or ResNet).

- It samples features from multiple layers: low-level layers provide fine-grained spatial detail, while high-level layers provide rich semantic context.
- By combining this multi-scale information for each pixel, PixelNet can make a highly informed and precise prediction without relying on upsampling layers, which can sometimes produce blurry or imprecise boundaries.
- This process is computationally intensive, so it is often applied to a sampled subset of pixels during training.

Applications: PixelNet is designed for tasks that require assigning a label to every single pixel in an image. Its main applications include:

- **Semantic Segmentation:** Labeling each pixel with class of object it belongs to (e.g., 'car', 'road', 'sky', 'person').
- **Object Part Segmentation:** Labeling the parts of an object (e.g., 'wheel', 'window', 'door' on a car).

- **Surface Normal Estimation:** Predicting orientation of surface at each pixel, which is useful in 3D reconstruction.

12. Define the following terms: Neuron, Weight, Bias, Activation Function.

- **Neuron:** A neuron (or node) is the fundamental computational unit of a neural network. It receives one or more inputs, performs a computation, and produces an output. The process involves:
 1. Calculating a **weighted sum** of its inputs.
 2. Adding a **bias** term to this sum.
 3. Passing the result through an **activation function** to produce the final output, which is then passed to other neurons.
- **Weight:** A weight is a learnable parameter associated with each input connection to a neuron. It represents the **strength or importance** of that connection. During training, the network learns the optimal values for these weights. A large positive weight means the input strongly excites the neuron, while a large negative weight means it strongly inhibits it.
- **Bias:** A bias is another learnable parameter associated with each neuron. It is added to the weighted sum of inputs before the activation function is applied. The role of the bias is to provide the neuron with a trainable constant that allows it to shift its activation function to the left or right. This helps the network fit the data better and is analogous to the y-intercept in a linear equation ($y=mx+b$, where b is the bias).
- **Activation Function:** An activation function is a mathematical function applied to the output of a neuron. Its purpose is to introduce **non-linearity** into the network. Without non-linear activation functions, a multi-layer neural network would simply be equivalent to a single-layer linear model, incapable of learning complex patterns. Common examples include:
 - **ReLU (Rectified Linear Unit):** $f(x)=\max(0,x)$
 - **Sigmoid:** $f(x)=1+e^{-x}$
 - **Tanh (Hyperbolic Tangent):** $f(x)=\tanh(x)$

13. Differentiate between feedforward neural network and recurrent neural network.

Feature	Feedforward Neural Network (FNN)	Recurrent Neural Network (RNN)
Data Flow	Unidirectional. Information flows in one direction only, from the input layer, through the hidden layers, to the output layer. There are no loops or cycles.	Cyclic. Information can flow in loops. The output of a neuron can be fed back to itself or to a neuron in a previous layer, creating a cycle.
Memory	No memory. FNNs have no internal memory of past inputs. The output for a given input is always the same, regardless of previous inputs.	Has internal memory. RNNs maintain a 'hidden state' or 'memory' that captures information from all previous inputs in a sequence.
Input/Output	Designed to handle fixed-size inputs and outputs.	Designed to handle sequential and variable-length data , such as text, speech, or time series.
Use Cases	Image classification, object detection, regression problems, and other tasks where the input data has no temporal dependency.	Natural Language Processing (NLP), speech recognition, machine translation, time-series forecasting, and other tasks involving sequential data.

Example Architectures	Multi-Layer Perceptron (MLP), Convolutional Neural Network (CNN).	Simple RNN, Long Short-Term Memory (LSTM), Gated Recurrent Unit (GRU).
------------------------------	--	--

14. What is the perceptron learning rule?

The **Perceptron Learning Rule** is a simple algorithm used to train a single-layer neural network, known as a Perceptron. It is an iterative algorithm designed to find a set of weights that allows the perceptron to correctly classify a set of linearly separable training data.

The Algorithm:

- 1. Initialization:** Initialize the weights (w) and bias (b) to small random numbers or zero.
- 2. Iteration:** For each training example (x_i, d_i) , where x_i is the input vector and d_i is the desired target label (e.g., +1 or -1):
 - a. Calculate the output:** Compute the neuron's output, y_i , by applying a step function to the weighted sum of inputs: $y_i = \text{step}(\sum w_j x_{ij} + b)$
(The step function typically outputs +1 if the input is > 0 , and -1 otherwise).
 - b. Compare and Update:** Compare the predicted output y_i with the desired output d_i .
 - If $y_i = d_i$ (prediction is correct), do nothing.
 - If $y_i \neq d_i$ (prediction is incorrect), update the weights and bias according to the following rule:
 $w_j(\text{new}) = w_j(\text{old}) + \eta(d_i - y_i)x_{ij}$
 $b(\text{new}) = b(\text{old}) + \eta(d_i - y_i)$
 Where η (eta) is the learning rate, a small positive constant that controls the step size of the update.
- 3. Termination:** Repeat step 2 for all training examples, iterating through the entire dataset multiple times (epochs) until the perceptron converges, meaning it classifies all training examples correctly.

Key Limitation: The Perceptron Learning Rule is only guaranteed to converge if the data is **linearly separable**.

15. State the challenges faced in multi-digit number recognition compared to single-digit.

Recognizing a sequence of multiple digits in an image is a significantly more complex problem than recognizing an isolated, single digit. The primary challenges include:

- 1. Segmentation:** Before the digits can be recognized, they must be located and separated from one another. In real-world images (like house numbers or handwritten sequences), digits can be touching, overlapping, or have inconsistent spacing, making it very difficult to determine where one digit ends and the next begins. Single-digit recognition tasks (like MNIST) typically provide pre-segmented, centered digits.
- 2. Variable Length:** The number of digits in the sequence is often unknown and variable (e.g., '42', '180', '90210'). The recognition system must be able to handle sequences of different lengths, which is a structural challenge for many models that expect a fixed-size input or output.
- 3. Spatial Context and Ordering:** The system must not only identify the digits but also preserve their correct order. Recognizing '1', '2', '3' is useless if the system cannot determine that the number is '123' and not '321' or '213'. The spatial relationship between the detected digits is crucial.
- 4. Increased Complexity:** The output space is exponentially larger. For single digits, there are 10 possible classes (0-9). For a five-digit number, there are $10^5 = 100,000$ possible classes. This makes the classification task vastly more complex.
- 5. Clutter and Noise:** Multi-digit numbers in natural scenes are often accompanied by background clutter, shadows, and varied lighting conditions that can interfere with both segmentation and recognition. The model needs to be robust enough to ignore these irrelevant details.

UNIT IV: Recurrent Neural Network (RNN)

1. Define sequence learning and give two examples of sequence-based tasks.

Sequence learning is a type of machine learning where the model works with data that is sequential in nature. In such data, the order of the elements is crucial for understanding its meaning and context. The goal is to learn a mapping from an input sequence to an output, which can be a single value or another sequence. Unlike traditional models that assume inputs are independent of each other, sequence learning models are designed to capture temporal dependencies and patterns over time.

Two examples of sequence-based tasks are:

1. **Machine Translation:** This is a classic sequence-to-sequence task. The input is a sequence of words in a source language (e.g., "How are you?"), and the model's output is a corresponding sequence of words in a target language (e.g., "Comment allez-vous ?"). The order of words is fundamental to the sentence's meaning.
2. **Sentiment Analysis:** This is often a sequence-to-one task. The input is a sequence of words, such as a movie review ("This movie was a fantastic and compelling story."), and the output is a single label that classifies the sentiment of the entire sequence (e.g., "Positive"). The model must understand the context built up by the sequence of words to make its final classification.

2. Differentiate between sequence-to-sequence and sequence-to-one models.

These models are differentiated by the nature of their output.

Feature	Sequence-to-One (Many-to-One)	Sequence-to-Sequence (Many-to-Many)
Input Type	A sequence of data points (e.g., words, time-series data).	A sequence of data points.
Output Type	A single, fixed-size output (e.g., a class label, a score).	A sequence of data points, which can be of a different length than the input.
Core Task	To classify or summarize an entire input sequence.	To transform an input sequence into an output sequence.
Typical Architecture	An RNN reads the entire input sequence and uses its final hidden state to make a prediction through a feedforward layer.	Typically uses an Encoder-Decoder architecture. The encoder processes the entire input sequence into a context vector, and the decoder generates the output sequence from that context.
Examples	- Sentiment Analysis (sequence of	- Machine Translation (sentence ->

	words -> positive/negative) - Video Classification (sequence of frames -> action label)	translated sentence) - Speech Recognition (audio sequence -> text sequence)
--	--	---

3. Why are RNNs suitable for sequence learning compared to feedforward networks?

RNNs are fundamentally better suited for sequence learning than feedforward networks (FNNs) due to their unique architectural features that address the limitations of FNNs.

1. **Internal Memory (Hidden State):** The most crucial feature of an RNN is its **hidden state**, which acts as a form of memory. At each time step, the RNN's output is a function of both the current input and the hidden state from the previous time step. This allows the network to retain information from past elements in the sequence and use it to process future elements. FNNs, in contrast, have no memory; they treat every input as an independent event.
2. **Parameter Sharing Across Time:** An RNN uses the same set of weights and biases for every time step in the sequence. This has two major benefits:
 - It allows the model to handle **variable-length sequences** gracefully, as the same mechanism is applied regardless of the sequence's length.
 - It enables the model to learn patterns that can occur at any point in the sequence. A pattern learned at the beginning of a sequence can be recognized if it appears at the end. FNNs would require separate weights for each position, which is inefficient and not scalable.
3. **Handling of Temporal Dependencies:** The recurrent loop structure inherently models the temporal flow of data, making RNNs naturally equipped to capture dependencies and relationships between elements that are far apart in the sequence. FNNs require a fixed-size input and lose all sequential ordering when data is flattened into a vector.

4. What is a computational graph? How is it used in training RNNs?

A **computational graph** is a directed graph used to represent a mathematical computation. In the context of neural networks, nodes in the graph represent variables (like inputs, weights) or operations (like matrix multiplication, addition, activation functions), and the edges represent the flow of data between these nodes.

How it's used in training RNNs:

The recurrent nature of an RNN, with its feedback loop, makes it difficult to apply standard backpropagation directly. To solve this, the RNN is "unrolled" (or "unfolded") in time to create a computational graph.

1. **Unrolling the Graph:** The compact, cyclic graph of an RNN is transformed into a deep, acyclic feedforward network structure. Each time step of the sequence becomes a distinct "layer" in this unrolled graph. The hidden state from time step $t-1$ is passed as an input to the layer for time step t . All these time-step layers share the same set of weights.
2. **Backpropagation Through Time (BPTT):** Once the RNN is unrolled, it behaves like a very deep feedforward network. This allows the standard backpropagation algorithm to be used for training. This specific application is called **Backpropagation Through Time (BPTT)**.
 - A forward pass is performed through the unrolled graph to compute the output and the loss at

each time step.

- A backward pass is then performed. Gradients of the total loss are calculated and propagated backward from the last time step to the first.
- Because the weights are shared across all time steps, the gradients for each weight matrix are summed up across all time steps before the weights are updated. This single update step incorporates the influence of the weight on the output across the entire sequence.

5. Differentiate between static and dynamic computational graphs.

Feature	Static Computational Graph	Dynamic Computational Graph
Definition	The graph structure is defined and compiled once before the model is run. The same graph is then executed repeatedly with different input data.	The graph structure is built on-the-fly as the code executes. The graph can be different for every iteration or input sample.
Flexibility	Rigid. The architecture is fixed. Handling variable-length inputs (like in RNNs) requires special techniques like padding and bucketing.	Highly Flexible. The graph can change dynamically based on the control flow of the programming language (e.g., loops, if-statements). This makes it very natural to handle variable-length sequences.
Debugging	More difficult. Errors often occur during graph compilation or execution, far from the code that defined the graph. Debugging requires specialized tools.	Easier. Since the graph is built during execution, you can use standard debuggers to set breakpoints and inspect values at any point in the forward pass.
Performance	Can be highly optimized by the framework before execution, as the entire graph is known. This often leads to better performance and memory efficiency.	May have slightly more overhead due to the repeated graph construction, but modern frameworks have largely minimized this gap.
Framework Examples	TensorFlow 1.x, Theano	PyTorch, TensorFlow 2.x (in Eager Mode)

6. Explain the architecture of a bidirectional RNN.

A **Bidirectional RNN (BiRNN)** is an architecture designed to capture context from both past and future elements in a sequence for any given point in time. A standard (unidirectional) RNN only considers past context, which can be a limitation for many tasks.

Architecture:

A BiRNN consists of two independent RNN layers that process the input sequence in opposite directions.

1. **Forward Layer:** This is a standard RNN layer that reads the input sequence from left to right (from

time step $t=1$ to $t=T$). It produces a sequence of forward hidden states: $(\overrightarrow{h}_1, \overrightarrow{h}_2, \dots, \overrightarrow{h}_T)$. Each state $\overrightarrow{h}_t$ summarizes the past information up to time t .

2. **Backward Layer:** This is another RNN layer that reads the input sequence from right to left (from time step $t=T$ to $t=1$). It produces a sequence of backward hidden states: $(\overleftarrow{h}_1, \overleftarrow{h}_2, \dots, \overleftarrow{h}_T)$. Each state $\overleftarrow{h}_t$ summarizes the future information from time t onwards.
3. **Output:** At each time step t , the final hidden state representation, h_t , is created by combining the outputs from both layers. This combination is typically done by concatenation, but summation or averaging can also be used.

$$h_t = [\overrightarrow{h}_t; \overleftarrow{h}_t]$$

This combined representation h_t provides a rich summary of the input at position t , incorporating information from its entire context.

7. Mention one advantage and one limitation of bidirectional RNNs.

- **Advantage: Complete Contextual Understanding**
The primary advantage of a BiRNN is its ability to access both past (preceding) and future (succeeding) information for every point in the sequence. This provides a much richer and more complete contextual representation, leading to significantly better performance on tasks like Named Entity Recognition (NER), machine translation, and text summarization, where the meaning of a word is often clarified by the words that follow it.
- **Limitation: Unsuitable for Real-Time Prediction**
The main limitation is that a BiRNN is not causal; it requires the entire input sequence to be available before it can make a prediction for any time step. This is because the backward layer must process the sequence from the very end. Therefore, BiRNNs cannot be used for real-time applications where predictions must be made online as data arrives, such as live speech transcription or stock price prediction.

9. What is the vanishing gradient problem in RNNs?

The **vanishing gradient problem** is a critical issue that occurs when training deep neural networks, and it is particularly severe in RNNs processing long sequences. It describes a situation where the gradients of the loss function with respect to the network's early-layer weights become exponentially small (i.e., they "vanish").

Mechanism in RNNs:

During Backpropagation Through Time (BPTT), the gradient is propagated backward from the end of the sequence to the beginning. To do this, it is repeatedly multiplied by the recurrent weight matrix (W_{hh}) at each time step.

If the values in this matrix are small (more formally, if its leading eigenvalue is less than 1), these repeated multiplications cause the gradient signal to shrink exponentially as it travels back through time.

Consequence:

When the gradient becomes vanishingly small, the weight updates for the earlier time steps are

negligible. This means the model cannot learn long-range dependencies—it struggles to connect events that are far apart in the sequence. For example, in the sentence "The writer of the books... is famous," the model would struggle to learn the grammatical agreement between "writer" and "is" because the gradient from "is" would vanish before it could reach and update the weights associated with processing "writer."

10. Explain exploding gradients with an example.

Exploding gradients is the opposite problem to vanishing gradients. It occurs when the gradients grow exponentially large, leading to massive, unstable weight updates during training. This can cause the model's weights to become so large that they are represented as NaN (Not a Number), effectively destroying the model and halting the training process.

Mechanism in RNNs:

Similar to the vanishing gradient problem, this occurs during BPTT due to the repeated multiplication of the gradient by the recurrent weight matrix (W_{hh}). If the values in this matrix are large (leading eigenvalue > 1), the gradient signal will grow exponentially as it is propagated backward through time.

Example:

Consider a very simple RNN with a single recurrent neuron where the update rule is $h_t = w_{hh} \cdot h_{t-1}$. When we compute the gradient of the loss at time t with respect to a much earlier hidden state h_k (where $k \ll t$), the chain rule dictates that the gradient will contain the term:

$$\frac{\partial h_t}{\partial h_k} = \prod_{i=k+1}^t \frac{\partial h_i}{\partial h_{i-1}} = \prod_{i=k+1}^t w_{hh} = (w_{hh})^{t-k}$$

If the absolute value of the weight $|w_{hh}|$ is greater than 1 (e.g., $w_{hh}=1.5$), this term will grow exponentially as the distance $(t-k)$ increases. A small error at the end of the sequence will be amplified into a massive gradient for the beginning of the sequence, causing the weights to "explode."

The most common technique to combat this is **gradient clipping**, where the gradients are manually scaled down if their norm exceeds a certain predefined threshold.

11. Draw and explain the basic architecture of an LSTM cell.

An **LSTM (Long Short-Term Memory)** cell is a sophisticated type of recurrent unit designed specifically to overcome the vanishing gradient problem and learn long-range dependencies. It achieves this using a system of gates to regulate the flow of information.

Core Components and Explanation:

The key to an LSTM is the **cell state (C_t)**, which acts as a memory "conveyor belt." Information can flow along it largely unchanged, making it easy for gradients to pass through many time steps. The flow is controlled by three main gates:

1. **Forget Gate (f_t)**: This gate decides what information to discard from the previous cell state (C_{t-1}). It looks at the previous hidden state (h_{t-1}) and the current input (x_t) and outputs a number between 0 and 1 for each number in the cell state. A 1 means "completely keep this," while a 0 means "completely forget this."

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

2. **Input Gate (i_t):** This gate decides what new information to store in the cell state. It has two parts:
 - A sigmoid layer (i_t) decides which values to update.
 - A tanh layer creates a vector of new candidate values ($\tilde{C}_t$) that could be added.

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

The old cell state is then updated by forgetting some information ($f_t \odot C_{t-1}$) and adding the new candidate information ($i_t \odot \tilde{C}_t$).

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$

3. **Output Gate (o_t):** This gate decides what part of the cell state will be output as the new hidden state (h_t). The cell state is passed through a tanh function (to squash values between -1 and 1) and then multiplied by the output of the sigmoid gate (o_t).

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t \odot \tanh(C_t)$$

This gating mechanism allows LSTMs to selectively remember or forget information over long periods, making them highly effective for complex sequence tasks.

13. Differentiate between LSTM and vanilla RNN.

Feature	Vanilla RNN (Simple RNN)	LSTM (Long Short-Term Memory)
Internal Structure	Very simple, consisting of a single recurrently connected layer, typically with a tanh or ReLU activation function.	Complex, containing an internal cell state and three regulatory gates (forget, input, and output).
Memory Handling	Has a single hidden state that serves as its short-term memory. This state is completely overwritten at each time step.	Maintains both a hidden state (short-term memory) and a cell state (long-term memory). The gates control what is added to or removed from this long-term memory.
Gradient Flow	Highly susceptible to the vanishing and exploding gradient problems due to repeated matrix multiplications in the backpropagation path.	The additive nature of the cell state update provides a clear path for gradients to flow, which significantly mitigates the vanishing gradient problem .
Capability	Effective at learning short-term dependencies but fails on tasks requiring the model to remember information over long sequences.	Specifically designed to learn long-range dependencies , making it highly effective for complex tasks like machine translation and speech recognition.

Computational Cost	Computationally cheaper and has fewer parameters.	More complex and computationally expensive due to the additional gating mechanisms and parameters.
---------------------------	---	--

14. How can RNNs be applied in self-driving cars?

Self-driving cars operate in a dynamic environment where understanding sequences of events over time is crucial for safety and decision-making. RNNs (and their variants like LSTMs and GRUs) are applied to model these temporal aspects.

1. **Trajectory Prediction:** One of the most critical tasks is predicting the future path of other agents (vehicles, pedestrians, cyclists). An RNN can take a sequence of past positions, velocities, and accelerations of an object as input and output a sequence of its most likely future coordinates. This is essential for anticipatory driving and collision avoidance.
2. **Sensor Fusion over Time:** A self-driving car is equipped with multiple sensors (cameras, LiDAR, radar). An RNN can be used to fuse the data streams from these sensors over time. By processing sequences of sensor readings, the model can build a more stable, robust, and noise-resistant representation of the surrounding environment, improving object detection and tracking.
3. **Behavior Prediction:** RNNs can be used to classify and predict the intent of other drivers or pedestrians. For example, by analyzing a sequence of a vehicle's movements (slight drift, turn signal status, speed changes), an RNN can predict the probability of an imminent lane change or turn.

15. Mention two sequence-learning problems faced by self-driving cars.

Self-driving cars must solve numerous complex sequence-learning problems to navigate safely. Two key examples are:

1. **Pedestrian Intent Prediction:** This is the problem of predicting a pedestrian's next action based on their recent behavior.
 - **Input:** A sequence of observations, such as video frames capturing their body posture, walking direction, speed, and head orientation over the last few seconds.
 - **Output:** A classification of their intent (e.g., 'will cross the street', 'will wait', 'is unaware of the car').
 - **Challenge:** Pedestrian behavior can be highly unpredictable and change abruptly. The model must learn subtle cues from the sequence to make a safe and timely decision. This is a critical sequence-to-one problem.
2. **Vehicle Trajectory Forecasting in Complex Scenarios:** This is the problem of predicting the future path of other vehicles, especially in dense traffic or at intersections.
 - **Input:** A sequence of the vehicle's past states (position, velocity, acceleration) and contextual information (e.g., traffic light status, road geometry).
 - **Output:** A sequence of future coordinates representing the vehicle's likely path over the next 3-5 seconds.
 - **Challenge:** A vehicle's future trajectory is influenced not just by its own history but also by the interactive behavior of all surrounding vehicles. The model must learn these complex multi-agent dynamics. This is a challenging sequence-to-sequence problem.

8. What is Backpropagation Through Time (BPTT)? Why is it needed for RNNs?

Backpropagation Through Time (BPTT) is the standard algorithm used to train Recurrent Neural Networks. It is an adaptation of the standard backpropagation algorithm applied to the "unrolled" version of an RNN.

Why is it needed?

A standard backpropagation algorithm works on feedforward networks, which have a clear, acyclic path from input to output. RNNs, however, contain cycles (or feedback loops) where the output of a time step is fed back as an input to the next. Standard backpropagation cannot handle these loops. BPTT solves this by creating a feedforward representation of the RNN, allowing gradients to be calculated.

The Process:

1. **Unroll the Network:** The first step is to unroll the RNN's recurrent structure through time. This transforms the network with a cycle into a very deep feedforward network, where each time step of the input sequence corresponds to one layer in the unrolled network. Critically, the **weights are shared** across all these time-step layers.
2. **Forward Pass:** An input sequence is passed through the unrolled network, and the outputs are calculated at each time step. The error (loss) is then computed by comparing the network's predictions to the true target sequence.
3. **Backward Pass:** The gradient of the loss is then calculated and propagated backward from the last time step to the first, just like in a standard deep feedforward network.
4. **Weight Update:** Because the weights (W_{xh} , W_{hh} , W_{hy}) are shared across all time steps, the gradients for each shared weight are summed up across all the unrolled steps. Finally, the weights are updated using this aggregated gradient.

A limitation of BPTT is that for very long sequences, the unrolled graph becomes extremely deep, which is computationally expensive and can lead to vanishing or exploding gradients. To manage this, a truncated version (Truncated BPTT) is often used, where backpropagation is only run for a fixed number of recent time steps.

12. Explain the architecture of a Gated Recurrent Unit (GRU).

A **Gated Recurrent Unit (GRU)** is a type of recurrent neural network cell, introduced as a simpler and more computationally efficient alternative to the LSTM. Like LSTMs, GRUs use gating mechanisms to control the flow of information and effectively combat the vanishing gradient problem, allowing them to capture long-range dependencies.

Architecture and Key Differences from LSTM:

A GRU simplifies the LSTM architecture in two main ways:

- It merges the cell state and the hidden state into a **single hidden state vector (h_t)**.
- It combines the LSTM's forget and input gates into a single **"update gate"**.

A GRU cell consists of two primary gates:

1. **Update Gate (z_t):** This gate determines how much of the past information (from the previous hidden state h_{t-1}) should be carried forward to the future. It acts like a combination of the forget

and input gates from an LSTM.

- It calculates a value between 0 and 1. A value close to 1 means the previous hidden state is almost entirely passed through, while a value close to 0 means the new candidate state is used instead.

$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

2. **Reset Gate (r_t):** This gate determines how much of the past information to *forget* when creating the new candidate hidden state. This allows the cell to drop information that is found to be irrelevant for the future.

- It also calculates a value between 0 and 1. If the value is close to 0, it effectively makes the cell act as if it's reading the first input of a new sequence, allowing it to forget the previously computed state.

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

These gates are then used to compute the new hidden state h_t . The reset gate decides how to combine the previous state with the current input to create a candidate state ($\tilde{h}_t$), and the update gate decides how to mix that candidate state with the original previous state (h_{t-1}) to get the final output.

Benefit: With fewer parameters and operations than an LSTM, a GRU is computationally faster and may perform better on smaller datasets where it is less likely to overfit.